

Exercise 7 - Linear and Logistic Regression using R

S.DEIVANAYKI (21BCS003)

12.03.2024


#Create a DataFrame with Columns as Price and Discount.

```
df<-  
data.frame(Price=c(368,340,665,954,331,556,376,332,788,826),Discount=c(1.7,1.  
5,2.8,5,1.3,2.2,1.3,1.2,3.75,3.9))  
head(df)
```

```
##   Price Discount  
## 1   368      1.7  
## 2   340      1.5  
## 3   665      2.8  
## 4   954      5.0  
## 5   331      1.3  
## 6   556      2.2
```

#Visualize the Dataframe


```
library(ggplot2)  
ggplot(df,aes(Price,Discount))+geom_point()
```



#1. Create models using geom_abline().

```
library(tidyverse)
```

```
models<-tibble(a1=runif(200,-1,1),a2=runif(200,-0.01,0.01))
ggplot(df,aes(Price,Discount))+geom_abline(aes(intercept=a1,slope=a2),data=models,alpha=0.3)+geom_point()
```




#Create model for specific slope and intercept

```
model1 <- function(a,data)
{
  a[1]+data$Price*a[2]
}
y_pred <- model1(c(-0.612,0.0056),df)
y_pred

## [1] 1.4488 1.2920 3.1120 4.7304 1.2416 2.5016 1.4936 1.2472 3.8008 4.0136
```

#Visualize Actual Vs Predicted Discount

```
ggplot(df,aes(Price))+geom_point(aes(y=Discount))+geom_line(aes(y=y_pred))
```



#Calculate Root mean square error

```
measure_distance <- function(mod,data)
{
  diff<-df$Discount-model1(mod,data)
  sqrt(mean(diff^2))
}
rmse <- measure_distance(c(-0.612,0.0056),df)
rmse

## [1] 0.2063883
```

#2.Generate intercept and slope using purrr


```
library(purrr)
xy_dist <- function(a1,a2)
{
  measure_distance(c(a1,a2),df)
}

models <- models %>% mutate(dist = purrr::map2_dbl(a1,a2,xy_dist))
head(models)

## # A tibble: 6 × 3
##       a1      a2  dist
##   <dbl>  <dbl> <dbl>
## 1 -0.144 -0.00266 4.49
## 2 -0.0110 0.00450 0.313
## 3 0.968 -0.00817 6.77
## 4 -0.298 0.00750 1.47
## 5 0.310 0.00324 0.666
## 6 0.252 0.00969 3.29
```


#Visualise the top 10 models

```
ggplot(df, aes(Price, Discount)) + geom_point(size=2, color="red") + geom_abline(aes(intercept=a1, slope=a2, color=-dist), data=filter(models, rank(dist) <= 10))
```



#Visualize the top 10 intercept and slope.

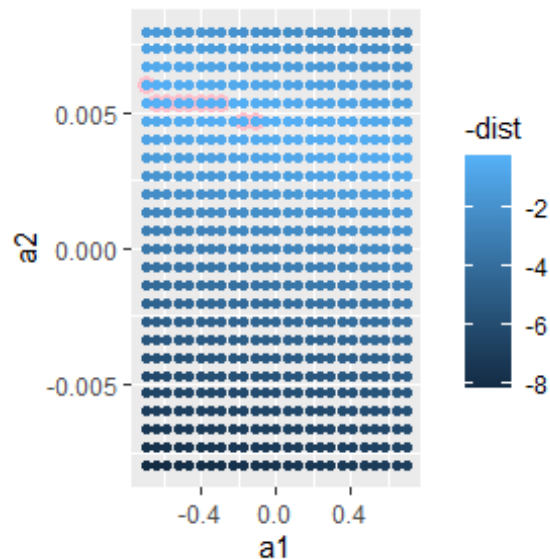
```
ggplot(models, aes(a1, a2)) + geom_point(data=filter(models, rank(dist) <= 10), size=3, color="blue") + geom_point(aes(colour=-dist))
```




#3. Visualize the Best Models using Grid Search.

```
grid <- expand.grid(a1=seq(-0.7, 0.7, length=25), a2=seq(-0.008, 0.008, length=25)) %>% mutate(dist = purrr::map2_dbl(a1, a2, xy_dist))
grid %>%
```

```
ggplot(aes(a1,a2))+geom_point(data=filter(grid,rank(dist)<=10),size=3,color="pink")+geom_point(aes(color=-dist))
```



```
ggplot(df,aes(Price,Discount)) + geom_point(size=2,color="blue") +  
geom_abline(aes(intercept=a1,slope=a2,color=-dist),data=filter(grid,rank(dist)<=10))
```




#4. Newton-Raphson Method to find slope and intercept

```
best <- optim(c(0,0),measure_distance,data=df)  
best$par
```

```
## [1] -0.603489635 0.005542682
```

```
ggplot(df,aes(Price,Discount)) + geom_point(size=2,color="blue") +  
geom_abline(intercept=best$par[1],slope=best$par[2])
```



#5. Create a linear model using `lm()`.

```
ind=sample(1:10,7)
data_train <- df[ind,]
data_test  <- df[-ind,]
linear_model <- lm(Discount~Price,data=data_train)
summary(linear_model)

##
## Call:
## lm(formula = Discount ~ Price, data = data_train)
##
## Residuals:
##      1      10      4      3      8      6      9
## 0.31129 -0.08133  0.29409 -0.26995  0.01507 -0.25293 -0.01623
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.6944397  0.3009045  -2.308   0.0691 .
## Price        0.0056607  0.0004443  12.740  5.3e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2559 on 5 degrees of freedom
## Multiple R-squared:  0.9701, Adjusted R-squared:  0.9641
## F-statistic: 162.3 on 1 and 5 DF,  p-value: 5.299e-05

coef(linear_model)

## (Intercept)      Price
## -0.694439675  0.005660744
```

```

y_pred<-predict(linear_model,data_test)
y_pred

##           2           5           7
## 1.230213 1.179266 1.434000

mse <- sum((y_pred-data_test$Discount)^2)
mse

## [1] 0.1053175

rmse <- mse^0.5
rmse

## [1] 0.3245266

```

#Create a Multivariate Linear Regression model.

```

mat<-sample(seq(60,100,5),50,replace = TRUE)
che<-sample(seq(60,100,5),50,replace = TRUE)
phy<-sample(seq(60,100,5),50,replace = TRUE)
rv <- sample(c(2.5,-3.8),50,replace=TRUE)
cutoff<-mat+che*0.5+phy*0.5
cutoff<-cutoff+rv
df2<-data.frame(Maths=mat,Chemistry=che,Physics=phy,Cutoff=cutoff)
head(df2)

##   Maths Chemistry Physics Cutoff
## 1    70         95      70  155.0
## 2    60         60      75  123.7
## 3    80         65      70  143.7
## 4    95         75      85  171.2
## 5    60        100      90  151.2
## 6    75        100      60  157.5

ind=sample(1:50,35)
data_train <- df2[ind,]
data_test  <- df2[-ind,]
linear_model <- lm(Cutoff ~ Maths+Chemistry+Physics,data=data_train)
summary(linear_model)

##
## Call:
## lm(formula = Cutoff ~ Maths + Chemistry + Physics, data = data_train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.763 -3.552  1.732  2.414  3.291
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -3.76790    5.58728  -0.674   0.505

```

```
## Maths      1.00237    0.04671   21.460 < 2e-16 ***
## Chemistry  0.53897    0.03996   13.487 1.64e-14 ***
## Physics    0.50768    0.04323   11.743 6.05e-13 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.184 on 31 degrees of freedom
## Multiple R-squared:  0.9665, Adjusted R-squared:  0.9632
## F-statistic: 297.7 on 3 and 31 DF,  p-value: < 2.2e-16

y_pred<-predict(linear_model,data_test)
y_pred

##          1          3          8         10         15         18         24         29
## 153.1380 146.9925 167.3913 154.4778 137.3468 154.9857 170.8034 193.3627
##          32          37          38          39          41          44          45
## 144.2977 180.3313 190.7329 144.9620 145.2749 152.7337 149.4659

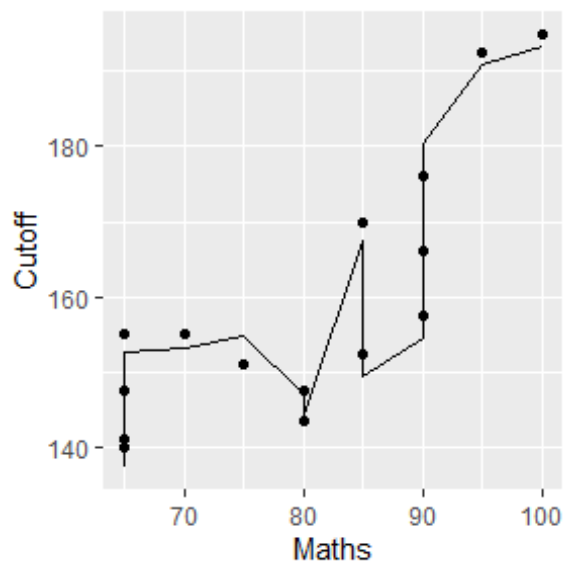
mse <- sum((y_pred-data_test$Cutoff)^2)
mse

## [1] 143.3228

rmse<-mse^0.5
rmse

## [1] 11.97175

ggplot(data_test,aes(Maths))+geom_point(aes(y=Cutoff))+geom_line(aes(y=y_pred
))
```



Common Question

#height (cms) vs age (years) for 120 persons by taking uniformly distributed values for age [10:70] and compute height using $h = a * R + 76$; where R is the random value in (6, 6.5, 7)

#1. Create the synthetic data.

```
age <- sample(10:70,120,replace=TRUE)
rv <- sample(c(6, 6.5, 7),120,replace=TRUE)
height <- 0.2*age*rv+76
df1<-data.frame(Age=age,Height=height)
head(df1)

##   Age Height
## 1  62  162.8
## 2  57  155.8
## 3  22  106.8
## 4  22  106.8
## 5  38  125.4
## 6  61  155.3
```

#2. Use simple linear regression and find the model.

```
ind=sample(1:120,84)
data_train <- df1[ind,]
data_test <- df1[-ind,]
linear_model <- lm(Height~Age,data=data_train)
summary(linear_model)

##
## Call:
## lm(formula = Height ~ Age, data = data_train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -7.5514 -2.9084 -0.1928  3.0494  6.2697
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  75.81313    1.15395   65.70  <2e-16 ***
## Age          1.31055    0.02584   50.71  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.743 on 82 degrees of freedom
## Multiple R-squared:  0.9691, Adjusted R-squared:  0.9687
## F-statistic: 2572 on 1 and 82 DF, p-value: < 2.2e-16

coef(linear_model)
```

```
## (Intercept)      Age
## 75.813133      1.310546

y_pred<-predict(linear_model,data_test)
y_pred[1:5]

##          5          7         13         19         20
## 125.6139 100.7135 164.9303 119.0612 104.6452
```

#3. Do performance evaluation for the model.

```
mse <- sum((y_pred-data_test$Height)^2)
mse


## [1] 388.0306

rmse <- mse^0.5
rmse

## [1] 19.69849
```

#4. Visualize the actual vs. predicted values.

```
ggplot(data_test,aes(Age))+geom_point(aes(y=Height))+geom_line(aes(y=y_pred))
```



#Logistic regression for the given California_housingdataset

```
library(MASS)

data <- housing
head(data)

##      Sat   Infl Type Cont Freq
## 1   Low   Low  Tower  Low   21
## 2 Medium   Low  Tower  Low   21
```

```
## 3   High   Low Tower   Low   28

data$Freq<-ifelse(data$Freq>25,1,0)
head(data)

##      Sat   Infl  Type Cont Freq
## 1    Low    Low Tower   Low    0
## 2 Medium    Low Tower   Low    0
## 3   High    Low Tower   Low    1

ind=sample(1:72,50)
data_train <- data[ind,]
data_test  <- data[-ind,]
log_model  <- glm(Freq~.,data=data_train,family = binomial)
summary(log_model)
## Call:
## glm(formula = Freq ~ ., family = binomial, data = data_train)
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    0.6628     1.4001   0.473   0.6359
## Sat.L          -0.7408     0.9397  -0.788   0.4305
## Sat.Q           2.4179     1.2176   1.986   0.0471 *
## InflMedium      1.1493     1.4116   0.814   0.4155
## InflHigh       -2.9915     1.5024  -1.991   0.0465 *
## TypeApartment   0.6535     1.2968   0.504   0.6143
## TypeAtrium     -22.9779    3656.8877  -0.006   0.9950
## TypeTerrace    -4.1870     1.8379  -2.278   0.0227 *
## ContHigh        1.6669     1.2014   1.387   0.1653
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 62.687  on 49  degrees of freedom
## Residual deviance: 24.886  on 41  degrees of freedom
## AIC: 42.886
##
## Number of Fisher Scoring iterations: 19

y_pred <- predict(log_model,data_test, type = "response")
y_pred

##              1              2              11              15              17
18
## 8.978659e-01 2.122471e-01 3.412058e-01 9.492561e-01 2.534611e-02
2.293563e-01
##              21              24              32              36              37
38
## 3.235094e-10 1.020995e-09 1.275366e-02 2.346619e-03 9.789707e-01
```

```

y_pred <- ifelse(y_pred > 0.5, 1, 0)
y_pred

## 1  2 11 15 17 18 21 24 32 36 37 38 40 41 42 43 46 49 55 69 70 71
## 1  0  0  1  0  0  0  0  0  0  1  1  1  1  1  1  1  1  0  0  0  0

mean(y_pred == data_test$Freq)

## [1] 0.6818182

```

#Program to predict diabetes Pima Indian Diabetes dataset using Logistic Regression Classifier.

```

data <- Pima.tr2
data<-data[rowSums(is.na(data)) == 0, ]
head(data)

##  npreg glu bp skin  bmi  ped age type
## 1     5  86 68   28 30.2 0.364 24  No
## 2     7 195 70   33 25.1 0.163 55  Yes
## 3     5  77 82   41 35.8 0.156 35  No

ind=sample(1:200,160)
data_train <- data[ind,]
data_test <- data[-ind,]
log_model <- glm(type~.,data=data_train,family = binomial)
summary(log_model)

## Call:
## glm(formula = type ~ ., family = binomial, data = data_train)
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -10.277630   1.999342  -5.141 2.74e-07 ***
## npreg        0.159694   0.071526   2.233  0.02557 *
## glu          0.033160   0.007599   4.364 1.28e-05 ***
## bp           0.003767   0.020921   0.180  0.85713
## skin        -0.011889   0.026353  -0.451  0.65187
## bmi          0.085330   0.048203   1.770  0.07669 .
## ped         2.116938   0.748305   2.829  0.00467 **
## age         0.028674   0.024987   1.148  0.25116
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 205.92  on 159  degrees of freedom
## Residual deviance: 140.18  on 152  degrees of freedom
## AIC: 156.18
##
## Number of Fisher Scoring iterations: 5

```

```

y_pred <- predict(log_model,data_test, type = "response")
y_pred <- ifelse(y_pred >0.5, "Yes", "No")
y_pred

##      3      6     11     16     20     24     25     32     33     36     38     42
45
## "No"   "No" "Yes"  "No"   "No"   "No"   "No"   "No"   "No" "Yes"  "No"   "No"
"No"
##    53    65    66    74    80    89    99   103   105   106   108   119
120
## "Yes"  "No"  "No"  "No"  "Yes"  "No"  "No"  "No"  "No"  "No"  "No"  "No"
"Yes"

mean(y_pred == data_test$type)

## [1] 0.75

```

#Program to predict whether the client will subscribe (1/0) to a term deposit (variable y).

```

library(readr)
data <- read_delim("C:/Users/DELL/Downloads/bank-additional/bank-
additional/bank-additional.csv",delim = ";")

data<-data[rowSums(is.na(data)) == 0, ]
data$y<-as.factor(data$y)
head(data)

## # A tibble: 6 × 21
##   age job      marital education default housing loan  contact month
day_of_week
##   <dbl> <chr>  <chr>    <chr>    <chr>    <chr>  <chr> <chr>  <chr> <chr>
## 1   30 blue-... married basic.9y no      yes    no    cellul... may  fri
## 2   39 servi... single  high.sch... no      no     no     teleph... may  fri
## 3   25 servi... married high.sch... no      yes    no     teleph... jun  wed
## 4   38 servi... married basic.9y no      unknown unkn... teleph... jun  fri
## 5   47 admin. married universi... no      yes    no     cellul... nov  mon
## 6   32 servi... single  universi... no      no     no     cellul... sep  thu
## # i 11 more variables: duration <dbl>, campaign <dbl>, pdays <dbl>,
## #   previous <dbl>, poutcome <chr>, emp.var.rate <dbl>, cons.price.idx
<dbl>,
## #   cons.conf.idx <dbl>, euribor3m <dbl>, nr.employed <dbl>, y <fct>

ind=sample(1:4119,2883)
data_train <- data[ind,]
data_test <- data[-ind,]
log_model <- glm(y~.,data=data_train,family = binomial)
summary(log_model)

##
## Call:
## glm(formula = y ~ ., family = binomial(link = "logit"), data = data_train)
##

```

```
## Coefficients: (1 not defined because of singularities)
##               Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -2.570e+02  1.574e+02  -1.633 0.102507
## age           4.560e-03  1.031e-02   0.442 0.658223
## jobstudent     1.262e-01  5.208e-01   0.242 0.808565
## jobtechnician  1.422e-01  2.893e-01   0.491 0.623175
## jobunemployed  5.369e-01  4.951e-01   1.084 0.278173
## jobunknown    -3.914e-01  9.365e-01  -0.418 0.675981
## maritalmarried 2.003e-01  3.071e-01   0.652 0.514302
## maritalsingle  1.946e-02  3.586e-01   0.054 0.956728
## educationprofessional.course -1.874e-01  4.165e-01  -0.450 0.652841
## defaultyes    -8.515e+00  5.354e+02  -0.016 0.987312
## housingunknown -5.590e-01  6.228e-01  -0.898 0.369446
## housingyes    -1.466e-01  1.732e-01  -0.846 0.397346
## loanunknown   NA         NA         NA      NA
## loanyes       -3.979e-02  2.346e-01  -0.170 0.865286
## contacttelephone -1.358e+00  3.823e-01  -3.552 0.000382 ***
## monthaug       9.784e-01  5.260e-01   1.860 0.062905 .
## monthdec       1.823e+00  8.706e-01   2.094 0.036290 *
## monthjul       9.695e-02  4.607e-01   0.210 0.833330
## monthsep       6.583e-01  7.452e-01   0.883 0.377014
## day_of_weekmon -3.749e-02  2.698e-01  -0.139 0.889486
## duration       5.942e-03  3.408e-04  17.434 < 2e-16 ***
## campaign      -1.382e-01  5.885e-02  -2.349 0.018840 *
## pdays        -4.119e-04  8.303e-04  -0.496 0.619858
## previous       7.962e-02  2.274e-01   0.350 0.726209
## poutcomenonexistent 7.199e-01  3.910e-01   1.841 0.065585 .
## poutcomesuccess 1.785e+00  8.235e-01   2.168 0.030160 *
## emp.var.rate  -1.139e+00  5.996e-01  -1.899 0.057550 .
## cons.price.idx 2.168e+00  1.047e+00   2.070 0.038463 *
## cons.conf.idx  7.342e-02  3.330e-02   2.204 0.027492 *
## euribor3m     -4.004e-01  5.225e-01  -0.766 0.443504
## nr.employed   1.038e-02  1.258e-02   0.825 0.409408
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## Null deviance: 2005.8 on 2882 degrees of freedom
## Residual deviance: 1034.8 on 2830 degrees of freedom
## AIC: 1140.8
##
## Number of Fisher Scoring iterations: 12

y_pred <- predict(log_model,data_test, type = "response")
y_pred <- ifelse(y_pred >0.5, "Yes", "No")
y_pred[1:10]

##      1      2      3      4      5      6      7      8      9     10
## "No"  "No"  "No" "Yes" "No"  "No"  "No"  "No"  "No"  "No"
```